



Introduction to Machine Learning

Gradient Descent and Normal Equations

Nikolay Manchev

July 30, 2016

London Machine Learning Study Group

Next events

<http://www.meetup.com/London-Machine-Learning-Study-Group>

Follow me

<https://twitter.com/nikolaymanchev>

Slides and code

Available at <https://github.com/nmanchev/MachineLearningStudyGroup>

Linear Regression and Gradient Descent

Normal Equations

Linear Regression and Gradient Descent

Univariate Linear Regression

Fitting a linear regression model

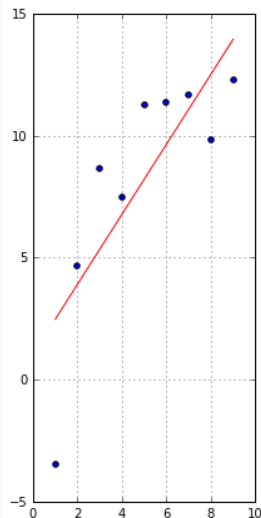
$$\mathbf{X} = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}^T$$

$$\mathbf{y} = \{y_1, y_2, \dots, y_N\}^T$$


Cost function minimisation

Minimising the cost function leads us to the coefficients of the best fitting line.

$$J(\mathbf{w}) = \frac{1}{2N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$$



Gradient Descent



Update rule

The derivative $\frac{d}{dw_0} J(w_0)$ provides the slope of the tangent line to the graph of the function at w_0

repeat until convergence {
 $w_0 := w_0 - \alpha \frac{d}{dw_0} J(w_0)$
}

- A positive $\alpha \frac{d}{dw_0} J(w_0)$ moves w_0 to the left
- A negative $\alpha \frac{d}{dw_0} J(w_0)$ moves w_0 to the right

Matrix Notation

$$\begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \vdots \\ \hat{y}_N \end{bmatrix} = \begin{bmatrix} 1 & x_{11} & x_{12} & \dots & x_{1D} \\ 1 & x_{21} & x_{22} & \dots & x_{2D} \\ \vdots & \vdots & \ddots & \vdots & \\ 1 & x_{N1} & x_{N2} & \dots & x_{ND} \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_D \end{bmatrix}$$

Hypothesis: $\hat{\mathbf{y}} = \mathbf{X}\mathbf{w}$

Cost function: $J(\mathbf{w}) = \frac{1}{2N} \sum_{i=1}^N (\mathbf{X}\mathbf{w} - \mathbf{y})^T (\mathbf{X}\mathbf{w} - \mathbf{y})$

Update rule:

repeat until convergence {

$$w_j := w_j - \alpha \frac{\mathbf{X}^T (\mathbf{X}\mathbf{w} - \mathbf{y})}{N}$$

Choice of alpha



Not an easy choice

$$w_0 := w_0 - \alpha \frac{d}{dw_0} J(w_0)$$

- Small alpha – slow convergence
- Large alpha – risk of overshooting

- Lipschitz continuity ($\alpha = \frac{1}{L}$)
- L not readily available – optimise manually
- Backtracking line search

Backtracking line search

Determine the maximum move along a given direction

Start with a large α and iteratively shrink it until an adequate decrease of the objective function.

Given a starting position \mathbf{w} and a search direction \mathbf{p} , we want to find a value of α that reduces $J(\mathbf{w} + \alpha \mathbf{p})$ relative to $J(\mathbf{w})$.

Backtracking Line Search

1. Select α_0 and control parameters $c \in (0, 1)$ and $\tau \in (0, 1)$
2. Set $j = 0$ and $t = -c \mathbf{p}$
3. **repeat** {
 $j := j + 1$
 $\alpha_j = \tau \alpha_{j-1}$
} **until** { $J(\mathbf{w}) - J(\mathbf{w} + \alpha_j \mathbf{p}) \geq \alpha_j t$ }

Generic Line Search

1. pick an initial w
2. **repeat until convergence** {
 - 2.1 calculate a search direction p (descent direction)
 - 2.2 find an optimal α
 - 2.3 $w := w + \alpha p$}

Normal Equations

Closed Form Solution

$$\frac{\partial}{\partial \mathbf{w}} J(\mathbf{w}) = \frac{\mathbf{X}^T(\mathbf{X}\mathbf{w} - \mathbf{y})}{N}$$

$$0 = \frac{\mathbf{X}^T(\mathbf{X}\mathbf{w} - \mathbf{y})}{N}$$

$$0 = \mathbf{X}^T(\mathbf{X}\mathbf{w} - \mathbf{y})$$

$$0 = \mathbf{X}^T\mathbf{X}\mathbf{w} - \mathbf{X}^T\mathbf{y}$$

$$\mathbf{X}^T\mathbf{X}\mathbf{w} = \mathbf{X}^T\mathbf{y}$$

$$\mathbf{w} = (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{y}$$

Gradient Descent vs. Normal Equations

Normal Equations

- No need to choose α
- No need to iterate to reach convergence
- Computing $(\mathbf{X}^T \mathbf{X})^{-1}$ is expensive – $O(n^3)$


Gradient Descent

- Works well for large n
- Can produce a “good enough” solution with a run-time that’s order of magnitude smaller [BB08]
- General optimisation algorithm

UCI Machine Learning Repository –

archive.ics.uci.edu/ml


- Great resource for Machine Learning data sets
- Over 330 freely available sets
- Auto MPG Data Set
 - Fuel consumption in MPG
 - Attributes: mpg, cylinders, displacement, horsepower, weight, acceleration etc.



Machine Learning Repository
Center for Machine Learning and Intelligent Systems

Auto MPG Data Set

Download: [Data Folder](#), [Data Set Description](#)




Abstract: Revised from CMU StatLib library, data concerns city-cycle fuel consumption

Data Set Characteristics:	Multivariate	Number of Instances:	398	Area:	N/A
Attribute Characteristics:	Categorical, Real	Number of Attributes:	8	Date Donated	1993-07-07
Associated Tasks:	Regression	Missing Values?	Yes	Number of Web Hits:	167833

MPG vs Horsepower

Simple use-case

- Auto MPG Data Set
- Predicting *MPG* based on *Horsepower*



Z-score Normalization

- Rescale the features to give them properties of a standard normal distribution ($\mu = 0, \sigma = 1$).
- Z-score normalization is calculated as

$$z = \frac{x - \mu}{\sigma}$$

where



μ is the mean of the population

σ is the standard deviation

- General requirement for many machine learning algorithms
- Helps Gradient Descent converge faster

Why normalizing the inputs works

- Gradient descent is curvature ignorant
- Brings all features to the same scale
- Gives the error surface a spherical shape. Look at [Hin14]

-  Olivier Bousquet and Léon Bottou, *The tradeoffs of large scale learning*, Advances in Neural Information Processing Systems 20 (J. C. Platt, D. Koller, Y. Singer, and S. T. Roweis, eds.), Curran Associates, Inc., 2008, pp. 161–168.
-  Geoffrey Hinton, *Neural Networks for Machine Learning, Lecture 6b*, Video Lecture, 2014.